

Regular expressions and automata

Adrien Mogenet

ContentSquare *am@csq.io*

September 9, 2016

- Regexes are simple (really? all the regexes?)

Objectives

- Regexes are simple (really? all the regexes?)
- Regexes are complex (no kidding!)

Objectives

- Regexes are simple (really? all the regexes?)
- Regexes are complex (no kidding!)
- Regexes can be fun (at least visual)

Objectives

- Regexes are simple (really? all the regexes?)
- Regexes are complex (no kidding!)
- Regexes can be fun (at least visual)
- There's an history behind the scene

Objectives

- Regexes are simple (really? all the regexes?)
- Regexes are complex (no kidding!)
- Regexes can be fun (at least visual)
- There's an history behind the scene
- Using regexes in your code is not harmless

Introducing Noam Chomsky



- American, born in 1928

Introducing Noam Chomsky



- American, born in 1928
- Linguist, philosopher, cognitive scientist, historian, logician, social critic, and political activist

Introducing Noam Chomsky



- American, born in 1928
- Linguist, philosopher, cognitive scientist, historian, logician, social critic, and political activist
- 50+ books in 50 years at MIT (Linguistics)

Introducing Noam Chomsky



- American, born in 1928
- Linguist, philosopher, cognitive scientist, historian, logician, social critic, and political activist
- 50+ books in 50 years at MIT (Linguistics)
- 100+ books (Politics)

Introducing Noam Chomsky



- American, born in 1928
- Linguist, philosopher, cognitive scientist, historian, logician, social critic, and political activist
- 50+ books in 50 years at MIT (Linguistics)
- 100+ books (Politics)
- 7 books (**ABOUT** Chomsky)

Introducing Noam Chomsky



- American, born in 1928
- Linguist, philosopher, cognitive scientist, historian, logician, social critic, and political activist
- 50+ books in 50 years at MIT (Linguistics)
- 100+ books (Politics)
- 7 books (**ABOUT** Chomsky)
- 20+ movies starring Chomsky

Introducing Noam Chomsky



- American, born in 1928
- Linguist, philosopher, cognitive scientist, historian, logician, social critic, and political activist
- 50+ books in 50 years at MIT (Linguistics)
- 100+ books (Politics)
- 7 books (**ABOUT** Chomsky)
- 20+ movies starring Chomsky
- Even got a bee named after him: *Megachile chomskyi* (2013)

Introducing Noam Chomsky



- American, born in 1928
- Linguist, philosopher, cognitive scientist, historian, logician, social critic, and political activist
- 50+ books in 50 years at MIT (Linguistics)
- 100+ books (Politics)
- 7 books (**ABOUT** Chomsky)
- 20+ movies starring Chomsky
- Even got a bee named after him: *Megachile chomskyi* (2013)
- Released his *Chomsky hierarchy* in 1950s

Introducing Noam Chomsky



- American, born in 1928
- Linguist, philosopher, cognitive scientist, historian, logician, social critic, and political activist
- 50+ books in 50 years at MIT (Linguistics)
- 100+ books (Politics)
- 7 books (**ABOUT** Chomsky)
- 20+ movies starring Chomsky
- Even got a bee named after him: *Megachile chomskyi* (2013)
- Released his *Chomsky hierarchy* in 1950s

- Hierarchy of classes of formal grammars

Formal grammars?

In an abstract world, set of production rules for a formal language.

- S : Start symbol
- A : Nonterminals (uppercase)
- a : Terminals (lowercase)
- ϵ : Empty string
- Σ : Alphabet (e.g. $\Sigma = a, b, c$)

Formal grammars?

Example of simple arithmetic language:

- $S \rightarrow \epsilon$
- $S \rightarrow A$
- $A \rightarrow n$ (any number)
- $A \rightarrow (A)$
- $A \rightarrow A + A$
- $A \rightarrow A - A$
- $A \rightarrow A * A$
- $A \rightarrow A / A$

Can generate something like $(2 + (3))/5$

Back to Chomsky hierarchy

Defined in 1956 by Noam Chomsky:

Grammar	Language	Production Rules
Type-0	Recursively enumerable (Turing machine)	$\alpha \rightarrow \beta$
Type-1	Context-sensitive	$\alpha A \beta \rightarrow \alpha \gamma \beta$
Type-2	Context-free	$A \rightarrow \gamma$
Type-3	Regular (Finite state automaton)	$A \rightarrow \alpha$ and $A \rightarrow \alpha A$

Chomsky hierarchy





Figure: <https://en.wikipedia.org/wiki/File:Chomsky-hierarchy.svg>


Without surprise, *regular expressions* are strongly linked to regular grammar (Type-3, generating regular languages).

Introducing Stephen Kleene




- American mathematician (1909 - 1994)

Introducing Stephen Kleene




- American mathematician (1909 - 1994)
- Student of Alonzo Church (along with Alan Turing, Emil Post)

Introducing Stephen Kleene




- American mathematician (1909 - 1994)
- Student of Alonzo Church (along with Alan Turing, Emil Post)
- got the National Medal of Science

Introducing Stephen Kleene




- American mathematician (1909 - 1994)
- Student of Alonzo Church (along with Alan Turing, Emil Post)
- got the National Medal of Science
- Formalized *regular languages*

Introducing Stephen Kleene



- American mathematician (1909 - 1994)
- Student of Alonzo Church (along with Alan Turing, Emil Post)
- got the National Medal of Science
- Formalized *regular languages*
- ...and invented regular expressions (1956)

Introducing Stephen Kleene



- American mathematician (1909 - 1994)
- Student of Alonzo Church (along with Alan Turing, Emil Post)
- got the National Medal of Science
- Formalized *regular languages*
- ...and invented regular expressions (1956)
- (Right after writing alternative proof to the Gödel's incompleteness theorems)

Regular language?

Definition

Regular language (or *rational language*) can be defined as a language recognized by a finite automaton. *i.e.* (equivalent properties):

- language of a regular expression

Regular language?

Definition

Regular language (or *rational language*) can be defined as a language recognized by a finite automaton. *i.e.* (equivalent properties):

- language of a regular expression
- accepted by a read-only Turing machine

Regular language?

Definition

Regular language (or *rational language*) can be defined as a language recognized by a finite automaton. *i.e.* (equivalent properties):

- language of a regular expression
- accepted by a read-only Turing machine
- accepted by a nondeterministic finite automaton (NFA)

Regular language?

Definition

Regular language (or *rational language*) can be defined as a language recognized by a finite automaton. *i.e.* (equivalent properties):

- language of a regular expression
- accepted by a read-only Turing machine
- accepted by a nondeterministic finite automaton (NFA)
- accepted by a deterministic finite automaton (DFA)

Regular language?

Formal definition (Production rules)

$$\begin{aligned} A &\rightarrow a \\ A &\rightarrow aB \end{aligned}$$

Examples

- Empty string language $\{\epsilon\} = \emptyset^*$ //
- Singleton language $\Sigma = \{a\}$ // aaaaaa
- Language over $\Sigma = \{a, b\}$ // aaaaabb

NOT a regular language

Type-1 grammar (context-sensitive)

$$\alpha A \beta \rightarrow \alpha \gamma \beta$$

NOT a regular language

Type-1 grammar (context-sensitive)

$$\alpha A \beta \rightarrow \alpha \gamma \beta$$

Famous example

- Set of strings $\{a^n b^n \mid n \geq 0\}$ // aaabbbb

NOT a regular language

Type-1 grammar (context-sensitive)

$$\alpha A \beta \rightarrow \alpha \gamma \beta$$

Famous example

- Set of strings $\{a^n b^n \mid n \geq 0\}$ // aaabbb

And so...

THIS CAN'T BE ACCEPTED BY A REGEX!

- French Wikipedia page title: “Expression rationnelle”

- French Wikipedia page title: “Expression rationnelle”
- Other available translation that sounds correct (to me): “Expression normale”

- French Wikipedia page title: “Expression rationnelle”
- Other available translation that sounds correct (to me): “Expression normale”
- These grammars and languages are everything but “régulier”

- French Wikipedia page title: “Expression rationnelle”
- Other available translation that sounds correct (to me): “Expression normale”
- These grammars and languages are everything but “régulier”
- Fun fact: other european countries chose equivalent of “régulier”:
Espressione regolare (it), Expressão regular (pt), Reguljära uttryck (sw), Regulārā izteiksme (le)

Please feed the troll

“Ah oui ? Qui d'autre que l'Office Québécois pour critiquer cette horreur de mot qui ressemble à l'anglais ?”

→ Moi, modeste informaticien et Français (...)

https://fr.wikipedia.org/wiki/Discussion:Expression_rationnelle

Are regexes really regular?

Meanwhile, in 2000s:

- Regexes exceed regular languages (Type-3) and became context-sensitive (Type-1):
- This: `(.+)\\1` \mapsto will match *papa* or *Pika!Pika!*.

Perl regexes performances

Exemple: recognize $a^n a^n$ against strings a^n .

```
time perl -e '("a" x 20) =~ /^(a?){20}(a){20}$/;'
real    0m0.132s
```

```
time perl -e '("a" x 26) =~ /^(a?){26}(a){26}$/;'
real    ????????
```

Perl regexes performances


Exemple: recognize $a^n a^n$ against strings a^n .

```
time perl -e '("a" x 20) =~ /^(a?){20}(a){20}$/;'
real    0m0.132s
```

```
time perl -e '("a" x 26) =~ /^(a?){26}(a){26}$/;'
real    0m11.507s
```

From regexes to automata

Regex = `kook`



Nondeterministic finite automaton

- Introduced in 1959

Nondeterministic finite automaton

- Introduced in 1959
- Recognize regular languages

Nondeterministic finite automaton

- Introduced in 1959
- Recognize regular languages
- Reading an input symbol is required for each state transition

Nondeterministic finite automaton

- Introduced in 1959
- Recognize regular languages
- Reading an input symbol is required for each state transition
- Its transitions are not uniquely determined by its source state and input symbol


Nondeterministic finite automaton

- Introduced in 1959
- Recognize regular languages
- Reading an input symbol is required for each state transition
- Its transitions are not uniquely determined by its source state and input symbol

And exist under many, many variations.

Regex? Automata?


Regex = `ko*k`



Another example

Regex = `abab|abbb`


Input = `abbb`



Another example

Regex = `abab|abbb`


Input = `abbb`



Another example

Regex = `abab|abbb`


Input = `a)bbb`



Another example

Regex = `abab|abbb`


Input = `ab⟩bb`



Another example

Regex = `abab|abbb`


Input = `abbb`



Another example

Regex = `abab|abbb`


Input = `a)bbb`



Another example

Regex = `abab|abbb`


Input = `ab⟩bb`



Another example

Regex = `abab|abbb`


Input = `abb)b`



Another example

Regex = `abab|abbb`

Input = `abbb`



Problem

Backtracking!

This is what happened with `(a?)26(a)26`. (At least) 2 alternatives:

- Thompson's NFA
- DFA (Deterministic Finite Automaton)

Alternative 1: Thompson's NFA

- Optimized by Ken Thompson, (Bell Telephone Labs Inc.), in a 1968 CACM paper
- Allows several possible states at the same moment



- American computer scientist (1943)



- American computer scientist (1943)
- Designed and implemented the original Unix operating system



- American computer scientist (1943)
- Designed and implemented the original Unix operating system
- Co-invented the Go programming language




- American computer scientist (1943)
- Designed and implemented the original Unix operating system
- Co-invented the Go programming language
- Turing Award (for *backdoor attack*)

Thompson's NFA example

Regex = `abab|abbb`


Input = `abbb`



Thompson's NFA example

Regex = `abab|abbb`


Input = `abbb`



Thompson's NFA example

Regex = `abab|abbb`


Input = `a)bbb`



Thompson's NFA example

Regex = `abab|abbb`


Input = `ab⟩bb`



Thompson's NFA example

Regex = `abab|abbb`


Input = `abb)b`



Thompson's NFA example

Regex = `abab|abbb`

Input = `abbb`






Figure: <https://swtch.com/~rsc/regexp/regexp1.html>

Alternative 2: DFA

Property

Each of its transitions is **uniquely** determined by its source state and input symbol


Theorem

Every NFA as an equivalent DFA [proof]

DFA example

Regex = `abab|abbb`


Input = `abbb`



DFA example

Regex = `abab|abbb`


Input = `a)bbb`



DFA example

Regex = `abab|abbb`


Input = `ab` \rangle `bb`



DFA example

Regex = `abab|abbb`


Input = `abb)b`



DFA example

Regex = `abab|abbb`

Input = `abbb`



Why don't everyone use Thompson's NFA or DFA???

You're not alone

- Java's `java.util.regex` uses backtracking too, so does PHP (PCRE library).
- Tcl, Awk, GNU Awk and GNU grep all use... DFA's!
- Google RE2 offers a C++ implementation (+ wrappers)

PCRE vs Google RE2

Regular expression	PCRE	RE2
Twain	5 ms	3 ms
(?i)Twain	79 ms	73 ms
[a-z]shing	564 ms	113 ms
Huck[a-zA-Z]+ Saw[a-zA-Z]+	30 ms	59 ms
\b\w+nn\b	837 ms	59 ms
[a-q][^u-z]{13}x	746 ms	3512 ms
Tom Sawyer Huckleberry Finn	40 ms	61 ms
(?i)Tom Sawyer Huckleberry Finn	424 ms	98 ms
.{0,2}(Tom Sawyer Huckleberry Finn)	5164 ms	66 ms
.{2,4}(Tom Sawyer Huckleberry Finn)	5298 ms	66 ms
Tom.{10,25}river river.{10,25}Tom	83 ms	68 ms
[a-zA-Z]+ing	1373 ms	129 ms
\s[a-zA-Z]{0,12}ing\s	592 ms	82 ms
([A-Za-z]awyer [A-Za-z]inn)\s	1112 ms	111 ms
["'][^"']{0,30}[?!\.]["']	65 ms	63 ms

Figure: Numbers from http://sljit.sourceforge.net/regex_perf.html

- Backreferences (e.g. `<(.)>(.*</(\1)>` to identify a closing XML tag)
- Space consumption of DFA can be way higher:
NFA: $O(n)$, DFA: $O(2^n)$, n is for regexp's length

Conclusion

- Be careful when using regexes
- Think about underlying engine
- Think about worst case usage (user's inputs...)
- Think about `strstr()`

- https://en.wikipedia.org/wiki/Regular_language
- https://en.wikipedia.org/wiki/Chomsky_hierarchy
- https://en.wikipedia.org/wiki/Stephen_Cole_Kleene
- <http://www.cs.ucr.edu/~jiang/cs215/tao-new.pdf>
- <https://www.debuggex.com/>
- <http://www.borntosegfault.com/2013/03/regexp-think-dfa.html>
- http://sljit.sourceforge.net/regex_perf.html
- https://en.wikipedia.org/wiki/Thompson's_construction

- <http://arstechnica.com/civis/viewtopic.php?f=20&t=1195549>
- http://www.perlmonks.org/?node_id=597262

- <https://en.wikipedia.org/wiki/File:Chomsky-hierarchy.svg>
- <https://en.wikipedia.org/wiki/File:Kleene.jpg>
- https://en.wikipedia.org/wiki/File:Noam_Chomsky_portrait_2015.jpg